

# Recommender System Case Study

Problem Statement: Sam's next exam would be to build a "Recommender System" using the Singular Value Decomposition (SVD) algorithm. Questions would be asked on the basis of what you've learnt in the respective module. Questions:

1. Implementing User-Based Recommender System using SVD (Singular Value Decomposition) method:
  - a. Load the 'ratings' and 'movies' datasets which is a part of 'MovieLens'
  - b. Find the unique number of users and movies in the 'ratings' dataset
  - c. Create a rating matrix for the 'ratings' dataset and store it in 'Ratings'
  - d. Load the 'ratings' dataset as SVD's Dataset object and compute 3-fold cross-validation using the SVD object
  - e. Find all the movies rated as 5 stars by user id '5' and store it in 'ratings\_1' data frame
  - f. Create a shallow copy of the 'movies' dataset and store the result in 'user\_5'
  - g. Train a recommender system using the SVD object and predict the ratings for user id '5'
  - h. Print the top10 movie recommendations for the user id '5'

```
In [1]: import os
os.chdir('C:\\Users\\veena\\OneDrive\\Desktop\\intellipaat assignment pdf s')
```

```
In [2]: import pandas as pd
ratings=pd.read_csv('ratings.csv')
ratings.head()
```

```
Out[2]:
```

	userId	movieId	rating	timestamp
0	1	2	3.5	1112486027
1	1	29	3.5	1112484676
2	1	32	3.5	1112484819
3	1	47	3.5	1112484727
4	1	50	3.5	1112484580

```
In [3]: movies=pd.read_csv("movies.csv")
movies.head()
```

```
Out[3]:
```

	movieId	title	genres
0	1	Toy Story (1995)	Adventure Animation Children Comedy Fantasy
1	2	Jumanji (1995)	Adventure Children Fantasy
2	3	Grumpier Old Men (1995)	Comedy Romance
3	4	Waiting to Exhale (1995)	Comedy Drama Romance
4	5	Father of the Bride Part II (1995)	Comedy

```
In [4]: n_users=ratings.userId.unique()
n_users
```

```
Out[4]: array([ 1, 2, 3, ..., 7118, 7119, 7120], dtype=int64)
```

```
In [5]: n_movies=ratings.movieId.unique()
n_movies
```

```
Out[5]: array([    2,    29,    32, ..., 32013, 102596, 65651], dtype=int64)
```

```
In [6]: print(f'number of users={n_users}and Number of movies={n_movies}')
```

```
number of users=[    1    2    3 ... 7118 7119 7120]and Number of movies=[    2
29    32 ... 32013 102596 65651]
```

```
In [7]: Ratings=ratings.pivot(index='userId',columns='movieId',values='rating').fillna(0)
Ratings.head()
```

```
Out[7]: movieId    1    2    3    4    5    6    7    8    9   10  ...  129350  129354  129428  129707  130
        userId
```

1	0.0	3.5	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	...	0.0	0.0	0.0	0.0
2	0.0	0.0	4.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	...	0.0	0.0	0.0	0.0
3	4.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	...	0.0	0.0	0.0	0.0
4	0.0	0.0	0.0	0.0	0.0	3.0	0.0	0.0	0.0	4.0	...	0.0	0.0	0.0	0.0
5	0.0	3.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	...	0.0	0.0	0.0	0.0

5 rows × 14026 columns

```
In [8]: # conda install -c conda-forge scikit-surprise
```

```
In [9]: from surprise import Reader, Dataset, SVD
from surprise.model_selection import cross_validate
```

```
In [10]: reader=Reader()
reader
```

```
Out[10]: <surprise.reader.Reader at 0x2475d117fa0>
```

```
In [11]: data=Dataset.load_from_df(ratings[['userId','movieId','rating']],reader)
```

```
In [12]: svd=SVD()
```

```
In [13]: cross_validate(svd,data,measures=['RMSE','MAE'],cv=3,verbose=True)
```

Evaluating RMSE, MAE of algorithm SVD on 3 split(s).

	Fold 1	Fold 2	Fold 3	Mean	Std
RMSE (testset)	0.8441	0.8459	0.8450	0.8450	0.0007
MAE (testset)	0.6469	0.6474	0.6479	0.6474	0.0004
Fit time	38.27	35.54	35.24	36.35	1.36
Test time	3.75	3.50	3.20	3.48	0.22

```
Out[13]: {'test_rmse': array([0.84413101, 0.8458596 , 0.84496193]),
'test_mae': array([0.64686998, 0.64735633, 0.64785583]),
'fit_time': (38.265305519104004, 35.5358145236969, 35.23546385765076),
'test_time': (3.7485849857330322, 3.4955708980560303, 3.2033848762512207)}
```

```
In [14]: ratings.head()
```

```
Out[14]:
```

	userId	movieId	rating	timestamp
0	1	2	3.5	1112486027
1	1	29	3.5	1112484676
2	1	32	3.5	1112484819
3	1	47	3.5	1112484727
4	1	50	3.5	1112484580

```
In [15]: ratings_1=ratings[(ratings['userId']==1)&(ratings['rating']==5)]
ratings_1=ratings_1.set_index('movieId')
ratings_1=ratings_1.join(movies)['title']
ratings_1.head()
```

```
Out[15]:
```

movieId	
4993	Honky Tonk Freeway (1981)
5952	Harder They Come, The (1973)
7153	Dreamers, The (2003)
8507	Barkleys of Broadway, The (1949)

Name: title, dtype: object

```
In [16]: user_1=movies.copy()
user_1=user_1.reset_index()
```

```
In [17]: data=Dataset.load_from_df(ratings[['userId','movieId','rating']],reader)
trainset=data.build_full_trainset()
svd.fit(trainset)
```

```
Out[17]: <surprise.prediction_algorithms.matrix_factorization.SVD at 0x2475cf15b50>
```

```
In [18]: user_1['Estimate_score']=user_1['movieId'].apply(lambda x:svd.predict(1,x).est)
user_1=user_1.drop(['movieId','genres','index'],axis=1)
user_1=user_1.sort_values('Estimate_score',ascending=False)
print(user_1.head(10))
```

	title	Estimate_score
5853	Lord of the Rings: The Two Towers, The (2002)	4.717112
7041	Lord of the Rings: The Return of the King, The...	4.704398
4897	Lord of the Rings: The Fellowship of the Ring,...	4.692378
10081	Taste of Tea, The (Cha no aji) (2004)	4.517837
657	Underground (1995)	4.501614
16897	Jane Eyre (2011)	4.499993
6666	Lagaan: Once Upon a Time in India (2001)	4.473225
15208	Cosmos (1980)	4.469709
7356	Band of Brothers (2001)	4.467460
7664	Gladiator (1992)	4.449940

```
In [ ]:
```

## Case Study 2 - K-means

Problem Statement: Consider yourself to be Sam who is a data scientist. He has been approached by a retail car showroom to help them segregate the cars into different clusters

Tasks to be performed:

1. Building the k-means clustering algorithm: a. Start off by extracting the 'mpg', 'disp' & 'hp' columns from the 'mtcars' data.frame. Store the result in 'car\_features' b. Build the kmeans algorithm on top of 'car\_features'. Here, the number of clusters should be 3 c. Bind the clustering vector to 'car\_features'. d. Extract observations belonging to individual clusters
2. On the same 'car\_features' dataset build a k-means algorithm, where the number of clusters is 5 a. Bind the clustering vector to 'car\_features' b. Extract observations belonging to individual clusters

```
In [ ]: import os
os.chdir('C:\\Users\\veena\\OneDrive\\Desktop')
```

```
In [20]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
%matplotlib inline
data=pd.read_csv('cars-3.csv')
data.head()
```

```
Out[20]:
```

	model	mpg	cyl	disp	hp	drat	wt	qsec	vs	am	gear	carb
0	Mazda RX4	21.0	6	160.0	110	3.90	2.620	16.46	0	1	4	4
1	Mazda RX4 Wag	21.0	6	160.0	110	3.90	2.875	17.02	0	1	4	4
2	Datsun 710	22.8	4	108.0	93	3.85	2.320	18.61	1	1	4	1
3	Hornet 4 Drive	21.4	6	258.0	110	3.08	3.215	19.44	1	0	3	1
4	Hornet Sportabout	18.7	8	360.0	175	3.15	3.440	17.02	0	0	3	2


```
In [21]: car_features=data[['mpg','disp','hp']]
car_features.head()
```

```
Out[21]:
```

	mpg	disp	hp
0	21.0	160.0	110
1	21.0	160.0	110
2	22.8	108.0	93
3	21.4	258.0	110
4	18.7	360.0	175

```
In [22]: f1=car_features['mpg'].values
f2=car_features['disp'].values
f3=car_features['hp'].values
x=np.array(list(zip(f1,f2,f3)))
plt.scatter(f1,f2,f3,c='blue')
```

```
Out[22]: <matplotlib.collections.PathCollection at 0x247a71a2700>
```



```
In [27]: import sklearn
```

```
In [29]: from sklearn.cluster import KMeans
```

```
In [33]: kmeans = KMeans(n_clusters=3)
kmeans = kmeans.fit(x)
labels = kmeans.predict(x)
centroids = kmeans.cluster_centers_
#print(C)
print(centroids)
```

C:\Users\veena\anaconda3\lib\site-packages\sklearn\cluster\\_kmeans.py:870: FutureWarning: The default value of `n\_init` will change from 10 to 'auto' in 1.4. Set the value of `n\_init` explicitly to suppress the warning

warnings.warn(  
C:\Users\veena\anaconda3\lib\site-packages\sklearn\cluster\\_kmeans.py:1382: UserWarning: KMeans is known to have a memory leak on Windows with MKL, when there are less chunks than available threads. You can avoid it by setting the environment variable OMP\_NUM\_THREADS=1.  
warnings.warn(  
[[ 14.64444444 388.22222222 232.11111111]  
[ 24.5 122.29375 96.875 ]  
[ 17.01428571 276.05714286 150.71428571]]

```
[[ 14.64444444 388.22222222 232.11111111]
 [ 24.5 122.29375 96.875 ]
 [ 17.01428571 276.05714286 150.71428571]]
```

```
In [ ]:
```

## KMeans Clustering Assignment

Problem Statement: You work in XYZ Company as a Python. The company officials want you to write code for a clustering problem. Dataset: customers.csv Tasks to be performed:

### 1. K-Means Clustering:

- Load customer data.
- Check the number of cells in each column with null values.
- Create a scatter plot with Age as X and Spending Score as Y.
- Draw a scatter plot displaying data points colored on the basis of clusters.

```
In [ ]:
```

```
In [1]: import os
os.chdir('C:\\Users\\veena\\OneDrive\\Desktop\\intellipaat assignment pdf s')
```

```
In [3]: import pandas as pd
import matplotlib.pyplot as plt
%matplotlib inline
```

```
In [5]: data=pd.read_csv('customers.csv')
data.head()
```

```
Out[5]:
```


	CustomerID	Gender	Age	Annual Income (k\$)	Spending Score (1-100)
0	1	Male	19	15	39
1	2	Male	21	15	81
2	3	Female	20	16	6
3	4	Female	23	16	77
4	5	Female	31	17	40

```
In [10]: data.isnull().sum()
```

```
Out[10]: CustomerID      0
Gender      0
Age      0
Annual Income (k$)      0
Spending Score (1-100)  0
dtype: int64
```

```
In [17]: x=data['Age']
y=data['Spending Score (1-100)']
plt.scatter(x,y)
```

```
Out[17]: <matplotlib.collections.PathCollection at 0x23899d6c370>
```



```
In [ ]:
```

## KMeans Clustering Assignment

Problem Statement: You work in XYZ Company as a Python Developer. The company officials want you to write code for a clustering problem. Dataset: customers.csv Tasks to be

performed:

### 1. K-Means Clustering:

- Load customer data.
- Check the number of cells in each column with null values.
- Create a scatter plot with Age as X and Spending Score as Y.
- Find out the best number for clusters between 1 and 10 (inclusive) using the elbowmethod.
- Draw a scatter plot displaying data points colored on the basis of clusters.

```
In [ ]: import os
os.chdir('C:\\Users\\veena\\OneDrive\\Desktop\\intellipaat assignment pdf s')
```

```
In [18]: import pandas as pd
import matplotlib.pyplot as plt
%matplotlib inline
```

```
In [19]: data=pd.read_csv('customers.csv')
data.head()
```

```
Out[19]:
```


	CustomerID	Gender	Age	Annual Income (k\$)	Spending Score (1-100)
0	1	Male	19	15	39
1	2	Male	21	15	81
2	3	Female	20	16	6
3	4	Female	23	16	77
4	5	Female	31	17	40

```
In [20]: data.isnull().sum()
```

```
Out[20]: CustomerID      0
Gender      0
Age      0
Annual Income (k$)      0
Spending Score (1-100)      0
dtype: int64
```

```
In [23]: x=data['Age']
y=data['Spending Score (1-100)']
plt.scatter(x,y)
```

```
Out[23]: <matplotlib.collections.PathCollection at 0x2389b4bd340>
```



```
In [24]: from sklearn.cluster import KMeans
```

```
In [47]: km=KMeans(n_clusters=3)
km
```

```
Out[47]: ▼      KMeans
          KMeans(n_clusters=3)
```

```
In [48]: y_predicted=km.fit_predict(data[['Age', 'Spending Score (1-100)']])
         y_predicted
```

```
C:\Users\veena\anaconda3\lib\site-packages\sklearn\cluster\_kmeans.py:870: FutureWarning: The default value of `n_init` will change from 10 to 'auto' in 1.4. Set the value of `n_init` explicitly to suppress the warning
  warnings.warn(
C:\Users\veena\anaconda3\lib\site-packages\sklearn\cluster\_kmeans.py:1382: UserWarning: KMeans is known to have a memory leak on Windows with MKL, when there are less chunks than available threads. You can avoid it by setting the environment variable OMP_NUM_THREADS=1.
  warnings.warn(
```

```
Out[48]: array([2, 1, 0, 1, 2, 1, 0, 1, 0, 1, 0, 1, 0, 1, 2, 1, 0, 1, 2, 1,
        0, 1, 0, 1, 2, 2, 0, 1, 0, 1, 0, 1, 0, 1, 0, 1, 2, 1, 2, 2,
        0, 1, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 1,
        2, 2, 1, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2,
        2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2,
        2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 1, 0, 1, 2, 1, 0, 1, 0, 1,
        2, 1, 0, 1, 0, 1, 0, 1, 0, 1, 2, 1, 0, 1, 2, 1, 0, 1, 0, 1,
        0, 1, 0, 1, 0, 1, 2, 1, 0, 1, 0, 1, 0, 1, 0, 1, 0, 1, 0, 1,
        0, 1, 0, 1, 2, 1, 0, 1, 2, 1, 0, 1, 0, 1, 0, 1, 0, 1, 0, 1,
        0, 1])
```

```
In [49]: data['cluster']=y_predicted
          data.head()
```


Out[49]:	CustomerID	Gender	Age	Annual Income (k\$)	Spending Score (1-100)	cluster
0	1	Male	19	15	39	2
1	2	Male	21	15	81	1
2	3	Female	20	16	6	0
3	4	Female	23	16	77	1
4	5	Female	31	17	40	2



```
In [50]: df1 = data[data.cluster==0]
df2 = data[data.cluster==1]
df3 = data[data.cluster==2]
plt.scatter(df1.Age,df1['Spending Score (1-100)'],color='green')
plt.scatter(df2.Age,df2['Spending Score (1-100)'],color='red')
plt.scatter(df3.Age,df3['Spending Score (1-100)'],color='yellow')
plt.xlabel('Age')
plt.ylabel('Spending Score (1-100)')
plt.legend()
```

No artists with labels found to put in legend. Note that artists whose label start with an underscore are ignored when legend() is called with no argument.

Out[50]: <matplotlib.legend.Legend at 0x2389c504550>



In [ ]:

## Agglomerative Clustering Assignment

Problem Statement: You work in XYZ Company as a Python Developer. The company officials want you to write code for an Agglomerative Clustering Problem. Tasks to be performed:

- Load iris data from load\_iris function from sklearn.datasets package.
- From the dataset extract the data property.
- Train an AgglomerativeClustering model based on the data.
- Plot dendrogram to visualize the clustering linkage

In [ ]:

```
In [ ]: import os
os.chdir('C:\\Users\\veena\\OneDrive\\Desktop\\intellipaath assignment pdf s')
```

```
In [3]: import pandas as pd
import matplotlib.pyplot as plt
%matplotlib inline
from sklearn import datasets
from sklearn.datasets import load_iris
from sklearn.cluster import AgglomerativeClustering
from scipy.cluster.hierarchy import dendrogram, linkage
```

```
In [4]: iris=datasets.load_iris()
```

```
In [5]: x_features=iris.data
        y_labels=iris.target
```


```
In [6]: model= AgglomerativeClustering(linkage='ward', n_clusters=3)
```

```
In [7]: model.fit(x_features)
        predicted_labels=model.labels_
```

```
In [8]: predicted_labels
```

```
Out[8]: array([1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1,
              1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1,
              1, 1, 1, 1, 1, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,
              0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 2, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,
              0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 2, 0, 2, 2, 2, 2, 0, 2, 2, 2,
              2, 2, 2, 0, 0, 2, 2, 2, 2, 0, 2, 0, 2, 0, 2, 2, 0, 0, 2, 2, 2, 2,
              2, 0, 0, 2, 2, 2, 0, 2, 2, 2, 0, 2, 2, 2, 0, 2, 2, 0], dtype=int64)
```

```
In [10]: linkage_matrix = linkage(x_features, 'ward')
         plot = plt.figure(figsize=(14,7))
         dendrogram(linkage_matrix, color_threshold=0,)
         plt.title('Dendrogram')
         plt.xlabel('sample index')
         plt.ylabel('distance')
         plt.show()
```



```
In [ ]:
```


## Dendrogram Assignment

Problem Statement: You work in XYZ Company as a Python. The company officials want you to write code for a Agglomerative Clustering Problem. Data:[[5,3], [10,15], [15,12], [24,10], [30,30], [85,70], [71,80], [60,78], [70,55], [80,91],] Tasks to be performed: Using the np.array function create an np array from the data given above. Generate a scatter plot for the data. Plot dendrogram to visualize the clustering linkage

```
In [11]: import numpy as np
         from scipy.cluster import hierarchy
         import matplotlib.pyplot as plt
```

```
In [15]: x = np.array([5.,3.,10.,15.,15.,12.,24.,10.,30.,
                       30.,85.,70.,71.,80.,60.])
```

```
In [16]: Z = hierarchy.linkage(x, 'single')
         plt.figure()
         dn = hierarchy.dendrogram(Z)
```



## Association Rule Mining Assignment

Problem Statement: You work in XYZ Company as a Python. The company officials want you to write code for a Association Rule Mining Dataset: retail\_dataset.csv Tasks to be performed:

- Using pandas import the dataset as dataframe
- Install the mixtend library to use apriori and association rule mining
- Using the apriori algorithm generate a list of item frequently brought together.
- Generate the association rules for the given items from apriori algorithm

```
In [1]: import os
         os.chdir('C:\\Users\\veena\\OneDrive\\Desktop\\intellipaath assignment pdf s')
```

```
In [2]: pip install apyori
```

Requirement already satisfied: apyori in c:\users\veena\anaconda3\lib\site-packages (1.1.2)  
Note: you may need to restart the kernel to use updated packages.

```
In [3]: import numpy as np
         import pandas as pd
         import matplotlib.pyplot as plt
         from apyori import apriori
```

```
In [6]: retail=pd.read_csv('C:\\Users\\veena\\OneDrive\\Desktop\\intellipaath assignment pdf s\\retail.csv')
         retail.head()
```

```
Input In [6]
retail=pd.read_csv("C:\Users\veena\OneDrive\Desktop\intellipaat assignment pdf s\retail.csv", sep='[,|:_]',encoding='unicode_escape',engine='python')
^
SyntaxError: (unicode error) 'unicodeescape' codec can't decode bytes in position
2-3: truncated \UXXXXXXXX escape
```

In [ ]: