

# Linear Regression Assignment

Problem Statement: You work in XYZ Company as a Python Data Scientist. The company officials have collected some data on salaries based on year of experience and wish for you to create a model from it. Dataset: data.csv Tasks to be performed:

1. Load the dataset using pandas
2. Extract data from YearsExperience column is a variable named X
3. Extract data from salary column is a variable named Y
4. Divide the dataset into two parts for training and testing in 66% and 33% proportion
5. Create and train LinearRegression Model on training set
6. Make predictions based on the testing set using the trained model
7. Check the performance by calculating the r2 score of the model

```
In [28]: import os
os.chdir('C:\\Users\\veena\\OneDrive\\Desktop\\intellipaath assignment pdf s')
```

```
In [29]: import pandas as pd
import matplotlib.pyplot as plt
%matplotlib inline
data=pd.read_csv('data.csv')
data.head()
```

```
Out[29]:
```

	YearsExperience	Salary
0	1.1	39343.0
1	1.3	46205.0
2	1.5	37731.0
3	2.0	43525.0
4	2.2	39891.0

```
In [30]: X=data[['YearsExperience']]
Y=data[['Salary']]
```

```
In [31]: from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
lr=LinearRegression()
```

```
In [32]: X_train, X_test, Y_train, Y_test = train_test_split(X,Y,test_size=0.33)
```

```
In [33]: lr.fit(X_train,Y_train)
```

```
Out[33]: LinearRegression()
```

```
In [34]: lr.predict(X_test)
```

```
Out[34]: array([[ 44568.93391277],
        [ 53109.2606687 ],
        [108146.9219847 ],
        [ 55956.03625401],
        [ 62598.51261974],
        [ 75883.46535119],
        [ 81577.01652181],
        [115738.32354553],
        [125227.57549657],
        [ 73985.61496098]])
```

```
In [35]: lr.score(X_test, Y_test)
```

```
Out[35]: 0.9791612020545786
```

## Logistic Regression Assignment

Problem Statement: You work in XYZ Company as a Python Data Scientist. The company officials have collected some data on Health parameter based on Diabetes and wish for you to create a model from it. Dataset: diabetes.csv Tasks to be performed:

1. Load the dataset using pandas
2. Extract data from Outcome column is a variable named Y
3. Extract data from every column except Outcome column in a variable named X
4. Divide the dataset into two parts for training and testing in 70% and 30% proportion
5. Create and train Logistic Regression Model on training set
6. Make predictions based on the testing set using the trained model
7. Check the performance by calculating the confusion matrix and accuracy score of the model

```
In [36]: df=pd.read_csv('diabetes.csv')
df.head()
```

```
Out[36]:
```

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	DiabetesPedigreeFunction	A
0	6	148	72	35	0	33.6		0.627
1	1	85	66	29	0	26.6		0.351
2	8	183	64	0	0	23.3		0.672
3	1	89	66	23	94	28.1		0.167
4	0	137	40	35	168	43.1		2.288

```
In [37]: X=df[['Pregnancies','Glucose','BloodPressure','SkinThickness','Insulin','BMI','DiabetesPedigreeFunction']]
Y=df[['Outcome']]
```

```
In [ ]: from sklearn.model_selection import train_test_split
```

```
In [38]: X_train, X_test, Y_train, Y_test = train_test_split(X,Y,test_size=0.3)
```

```
In [39]: from sklearn.linear_model import LogisticRegression
```

```
In [40]: lr=LogisticRegression()
```

```
In [41]: lr.fit(X_train,Y_train)
```

C:\Users\veena\anaconda3\lib\site-packages\sklearn\utils\validation.py:993: DataConversionWarning: A column-vector y was passed when a 1d array was expected. Please change the shape of y to (n\_samples, ), for example using ravel().

```
y = column_or_1d(y, warn=True)
```

C:\Users\veena\anaconda3\lib\site-packages\sklearn\linear\_model\\_logistic.py:814:

ConvergenceWarning: lbfgs failed to converge (status=1):

STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

Increase the number of iterations (max\_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

[https://scikit-learn.org/stable/modules/linear\\_model.html#logistic-regression](https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression)

```
n_iter_i = _check_optimize_result(
```

```
LogisticRegression())
```

```
In [42]: lr.predict(X_test)
```

```
Out[42]: array([1, 0, 1, 0, 0, 0, 1, 0, 0, 0, 1, 0, 0, 1, 1, 0, 0, 0, 0, 1, 0, 1,
        0, 0, 0, 0, 1, 1, 0, 0, 0, 1, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 1, 0,
        0, 1, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 1, 0, 0, 0, 1, 0, 0,
        0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 1, 0, 0, 1, 1, 0, 0, 1,
        0, 0, 0, 0, 0, 1, 0, 0, 0, 1, 1, 0, 0, 1, 1, 0, 0, 0, 0, 0, 0, 0,
        0, 0, 0, 0, 1, 0, 1, 1, 0, 1, 0, 0, 0, 1, 1, 0, 0, 1, 1, 1, 0, 1,
        0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 1, 1, 0, 1, 0, 1, 0, 0, 0,
        1, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0,
        0, 0, 0, 1, 0, 0, 1, 1, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,
        0, 1, 0, 1, 1, 0, 0, 1, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0,
        0, 1, 0, 0, 1, 0, 0, 0, 1, 0, 0], dtype=int64)
```

```
In [34]: lr.score(X_test,Y_test)
```

```
Out[34]: 0.8008658008658008
```

## Decision Tree Assignment

Problem Statement: You work in XYZ Company as a Python Data Scientist. The company officials have collected some data on salaries based on year of experience and wish for you to create a model from it. Dataset: diabetes.csv Tasks to be performed:

1. Load the dataset using pandas
2. Extract data from Outcome column is a variable named Y
3. Extract data from every column except Outcome column in a variable named X
4. Divide the dataset into two parts for training and testing in 70% and 30% proportion
5. Create and train Decision Tree Model on training set
6. Make predictions based on the testing set using the trained model
7. Check the performance by calculating the confusion matrix and accuracy score of the model

```
In [43]: df1=pd.read_csv('diabetes.csv')
df1.head()
```

Out[43]:	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	DiabetesPedigreeFunction	A
0	6	148	72	35	0	33.6		0.627
1	1	85	66	29	0	26.6		0.351
2	8	183	64	0	0	23.3		0.672
3	1	89	66	23	94	28.1		0.167
4	0	137	40	35	168	43.1		2.288

```
In [44]: X=df1[['Pregnancies', 'Glucose', 'BloodPressure', 'SkinThickness', 'Insulin', 'BMI', 'Diabetes', 'Outcome']]
```

```
In [45]: X_train, X_test, Y_train, Y_test = train_test_split(X,Y,test_size=0.3)
```

```
In [46]: from sklearn import tree
```

```
In [47]: tr=tree.DecisionTreeClassifier()
```

```
In [48]: tr.fit(X_test, Y_test)
```

```
Out[48]: DecisionTreeClassifier()
```

```
In [49]: tr.predict(X_test)
```

```
Out[49]: array([1, 0, 0, 1, 1, 1, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0,
1, 0, 1, 0, 1, 0, 1, 0, 0, 0, 0, 0, 0, 0, 1, 1, 0, 0, 0, 1,
1, 0, 0, 1, 0, 1, 1, 0, 0, 0, 0, 0, 0, 0, 0, 1, 1, 1, 0, 1, 1,
1, 0, 1, 0, 0, 0, 0, 0, 1, 1, 0, 0, 0, 0, 0, 1, 0, 0, 0, 1, 0, 0,
1, 0, 0, 0, 1, 0, 0, 0, 1, 1, 0, 1, 0, 0, 1, 0, 0, 0, 0, 1, 0, 1,
0, 0, 0, 1, 1, 0, 1, 1, 1, 0, 1, 0, 0, 1, 0, 1, 0, 1, 1, 0, 1, 1,
1, 1, 0, 0, 1, 0, 1, 1, 0, 1, 1, 1, 1, 0, 0, 1, 0, 0, 0, 0, 0, 0,
0, 0, 0, 0, 1, 1, 0, 0, 0, 0, 0, 0, 1, 0, 1, 0, 0, 1, 0, 1, 0, 1,
0, 0, 0, 0, 1, 0, 0, 0, 0, 1, 0, 0, 1, 1, 0, 1, 1, 1, 0, 0, 1, 0,
1, 1, 1, 1, 0, 1, 1, 0, 0, 0, 0, 1, 0, 1, 1, 0, 1, 0, 1, 0, 0, 1,
0, 1, 0, 0, 1, 1, 1, 1, 0, 1, 0], dtype=int64)
```

```
In [50]: tr.score(X_test, Y_test)
```

```
Out[50]: 1.0
```

```
In [51]: y_pred=tr.predict(X_test)
```

```
In [52]: from sklearn.metrics import confusion_matrix
cm=confusion_matrix(Y_test,y_pred)
cm
```

```
Out[52]: array([[140,  0],
                [ 0, 91]], dtype=int64)
```

In [ ]:

# Random Forest Assignment

Problem Statement: You work in XYZ Company as a Python Data Scientist. The company officials have collected some data on salaries based on year of experience and wish for you to create a model from it. Dataset: diabetes.csv Tasks to be performed:

1. Load the dataset using pandas
2. Extract data from Outcome column is a variable named Y
3. Extract data from every column except Outcome column in a variable named X.
4. Divide the dataset into two parts for training and testing in 70% and 30% proportion
5. Create and train Random Forest Model on training set
6. Make predictions based on the testing set using the trained model
7. Check the performance by calculating the confusion matrix and accuracy score of the model

```
In [1]: import os
os.chdir('C:\\Users\\veena\\OneDrive\\Desktop\\intellipaath assignment pdf s')
```

```
In [2]: import pandas as pd
import matplotlib.pyplot as plt
%matplotlib inline
```

```
In [3]: df2=pd.read_csv('diabetes.csv')
df2.head()
```

```
Out[3]:
```

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	DiabetesPedigreeFunction	Outcome
0	6	148	72	35	0	33.6	0.627	0
1	1	85	66	29	0	26.6	0.351	0
2	8	183	64	0	0	23.3	0.672	0
3	1	89	66	23	94	28.1	0.167	0
4	0	137	40	35	168	43.1	2.288	0

```
In [5]: df2.isnull().sum()
```

```
Out[5]: Pregnancies      0
Glucose      0
BloodPressure  0
SkinThickness  0
Insulin      0
BMI          0
DiabetesPedigreeFunction  0
Age          0
Outcome      0
dtype: int64
```

```
In [4]: a=df2[['Pregnancies','Glucose','BloodPressure','SkinThickness','Insulin','BMI','DiabetesPedigreeFunction']]
b=df2[['Outcome']]
```

```
In [7]: from sklearn.model_selection import train_test_split
```

```
In [8]: a_train, b_train, a_test, b_test = train_test_split(a,b,test_size=0.3)
```

```
In [9]: from sklearn.ensemble import RandomForestClassifier
rfc=RandomForestClassifier()
```

rfc

Out[9]: ▾ RandomForestClassifier  
RandomForestClassifier()

In [13]: df2.shape

Out[13]: (768, 9)

In [11]: rfc.fit(a\_train,b\_train)

```
-----
ValueError                                Traceback (most recent call last)
Input In [11], in <cell line: 1>()
----> 1 rfc.fit(a_train,b_train)

File ~\anaconda3\lib\site-packages\sklearn\ensemble\_forest.py:346, in BaseForest.fit(self, X, y, sample_weight)
    344 if issparse(y):
    345     raise ValueError("sparse multilabel-indicator for y is not supported.")
--> 346 X, y = self._validate_data(
    347     X, y, multi_output=True, accept_sparse="csc", dtype=DTYPE
    348 )
    349 if sample_weight is not None:
    350     sample_weight = _check_sample_weight(sample_weight, X)

File ~\anaconda3\lib\site-packages\sklearn\base.py:554, in BaseEstimator._validate_data(self, X, y, reset, validate_separately, **check_params)
    552 y = check_array(y, input_name="y", **check_y_params)
    553 else:
--> 554 X, y = check_X_y(X, y, **check_params)
    555 out = X, y
    557 if not no_val_X and check_params.get("ensure_2d", True):

File ~\anaconda3\lib\site-packages\sklearn\utils\validation.py:1122, in check_X_y(X, y, accept_sparse, accept_large_sparse, dtype, order, copy, force_all_finite, ensure_2d, allow_nd, multi_output, ensure_min_samples, ensure_min_features, y_numeric, estimator)
    1104 X = check_array(
    1105     X,
    1106     accept_sparse=accept_sparse,
    (...)
    1117     input_name="X",
    1118 )
    1120 y = _check_y(y, multi_output=multi_output, y_numeric=y_numeric, estimator=estimator)
-> 1122 check_consistent_length(X, y)
    1124 return X, y

File ~\anaconda3\lib\site-packages\sklearn\utils\validation.py:397, in check_consistent_length(*arrays)
    395 uniques = np.unique(lengths)
    396 if len(uniques) > 1:
--> 397     raise ValueError(
    398         "Found input variables with inconsistent numbers of samples: %r"
    399         % [int(1) for l in lengths]
    400     )

ValueError: Found input variables with inconsistent numbers of samples: [537, 231]
```

In [ ]:

# Naïve-Bayes Assignment

Problem Statement: You work in XYZ Company as a Python Data Scientist. The company officials have collected some data on salaries based on year of experience and wish for you to create a model from it. Dataset: diabetes.csv Tasks to be performed:

1. Load the dataset using pandas
2. Extract data from Outcome column is a variable named Y
3. Extract data from every column except Outcome column in a variable named X
4. Divide the dataset into two parts for training and testing in 70% and 30% proportion
5. Create and train Naïve Bayes Model on training set
6. Make predictions based on the testing set using the trained model
7. Check the performance by calculating the confusion matrix and accuracy score of the model

In [ ]:

```
In [3]: import os
os.chdir('C:\\Users\\veena\\OneDrive\\Desktop\\intellipaath assignment pdf s')
```

```
In [4]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
df3=pd.read_csv('diabetes.csv')
df3.head()
```

```
Out[4]:
```

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	DiabetesPedigreeFunction	A
0	6	148	72	35	0	33.6	0.627	
1	1	85	66	29	0	26.6	0.351	
2	8	183	64	0	0	23.3	0.672	
3	1	89	66	23	94	28.1	0.167	
4	0	137	40	35	168	43.1	2.288	

```
In [5]: x= df3.drop('Outcome',axis='columns')
y=df3.Outcome
```

```
In [7]: x.head()
```

Out[7]:	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	DiabetesPedigreeFunction	A
0	6	148	72	35	0	33.6		0.627
1	1	85	66	29	0	26.6		0.351
2	8	183	64	0	0	23.3		0.672
3	1	89	66	23	94	28.1		0.167
4	0	137	40	35	168	43.1		2.288

In [8]: `y.head()`

Out[8]:

0	1
1	0
2	1
3	0
4	1

Name: Outcome, dtype: int64

In [9]: `from sklearn.model_selection import train_test_split`

In [10]: `x_train, x_test,y_train,y_test=train_test_split(x,y,test_size=0.3,random_state=0)`

In [11]: `from sklearn.preprocessing import StandardScaler  
sc=StandardScaler()`

In [12]: `x_train=sc.fit_transform(x_train)  
x_test=sc.fit_transform(x_test)`

In [13]: `from sklearn.naive_bayes import GaussianNB  
gsb=GaussianNB()  
gsb.fit(x_train, y_train)`

Out[13]: GaussianNB()

In [21]: `y_pred=gsb.predict(x_test)`

In [24]: `gsb.predict(x_test)`

Out[24]:

```
array([1, 0, 0, 1, 0, 0, 1, 1, 1, 1, 1, 1, 0, 1, 0, 0, 1, 0, 0, 0, 1, 0,
       0, 0, 0, 0, 1, 0, 0, 1, 0, 0, 0, 0, 1, 1, 0, 0, 1, 0, 0, 0, 1,
       1, 0, 0, 0, 0, 0, 0, 1, 1, 0, 0, 0, 0, 0, 1, 1, 0, 0, 1, 1,
       1, 0, 0, 0, 0, 0, 1, 1, 0, 1, 1, 0, 0, 1, 0, 0, 0, 0, 0,
       1, 0, 0, 0, 0, 1, 0, 0, 1, 1, 0, 1, 0, 0, 1, 0, 0, 0, 0, 1,
       0, 1, 1, 1, 1, 1, 0, 1, 0, 1, 0, 1, 0, 1, 0, 0, 0, 0, 0, 0,
       0, 1, 0, 1, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,
       1, 0, 1, 0, 0, 1, 1, 1, 0, 0, 1, 0, 1, 0, 0, 0, 0, 0, 0, 1, 0,
       0, 0, 0, 0, 1, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 1, 1,
       0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0,
       0, 0, 1, 1, 1, 1, 0, 0, 0, 0, 1], dtype=int64)
```

In [23]: `gsb.score(x_test,y_test)`

Out[23]: 0.7272727272727273

In [26]: `from sklearn.metrics import confusion_matrix  
cm=confusion_matrix(y_test,y_pred)  
cm`



```
Out[26]: array([[128, 29],  
               [ 34, 40]], dtype=int64)
```

```
In [ ]:
```